

Homework 3 (final project)

```
In [1]: using PyPlot
using LinearAlgebra
using Statistics
using DelimitedFiles
using Distributions
using KernelDensity

# Univariate simulators

Gaussian(mu::Real, Theta::Real) = mu + sqrt(Theta) \ randn()

function Wishart(phi::Real, Phi::Real)
    d = 1
    n = Int(2phi)
    @assert n == 2phi "first argument must be Integer/2"
    @assert phi > d/2 - 1/2 "first argument must be > d/2 - 1/2"
    (2Phi) \ sum(randn(n).^2)
end

# Multivariate simulators

function Gaussian_d(mu::Vector, Theta::Symmetric)
    d = size(Theta,1)
    Theta1/2 = sqrt(Theta)
    mu + Theta1/2 \ randn(d)
end

function Wishart_d(phi::Real, Phi::Symmetric)
    n = Int(2phi)
    d = size(Phi,1)
    @assert n == 2phi "first argument must be Integer/2"
    @assert phi > d/2 - 1/2 "first argument must be > d/2 - 1/2"
    iW = zeros(d,d)
    zi = zeros(d)
    for i=1:n
        zi .= randn(d)
        iW .+= zi * zi'
    end
    twPhi1/2 = sqrt(Symmetric(2Phi))
    return Symmetric(twPhi1/2 \ iW / twPhi1/2)
end

# Misc
```

```
" 1130

function viridis_colors(;length=10)
  array_rgba = plt.cm.get_cmap("viridis", length)(range(0,1,length=length))
  return collect(eachrow(array_rgba))
end

function nearest_valid_phi(phi;d=1)
  lower_bound = d/2 - 1/2
  new_phi = round(2*max(phi, lower_bound)) / 2
  if new_phi != phi
    @warn "Input phi = $phi, nearest valid phi = $new_phi"
  end
  return new_phi
end
```

Out[1]: nearest_valid_phi (generic function with 1 method)

Data

The data for this problem is in the file `data.csv` . Here is how to load it.

```
In [2]: yi = readlm("data.csv", ',')[:]
```

```
Out[2]: 25-element Array{Float64,1}:
-2017.7106233488373
-2007.2408980724263
-2005.1154014930028
-2008.5117830345864
-2025.5132574522222
-2028.1854456826702
-2010.5135720350138
-2025.3244890587323
-2018.6094268648226
-2020.6825199021844
-2008.568246172282
-2011.299706617491
-2043.3103043391645
-2017.8271270156313
-2027.1828129414023
-1996.1439726063504
-2037.4334342224986
-2024.9749373753116
-2021.822162447521
-2024.1926752522218
-2013.991961096987
-2020.5517441255515
-2034.3415531928802
-2006.4170649918367
-2041.705257535142
```

I simulated this data by first choosing two numbers μ and θ . Then I simulated 25 numbers as follows

$$y_1, y_2, \dots, y_{25} \stackrel{iid}{\sim} \text{Gaussian}(\mu, \theta)$$

Note: as we did in the lecture notes, θ represents inverse variance.

The goal of this homework is to try to infer the numbers μ and $1/\theta$ using a Gibbs chain to sample from the posterior $p(\mu, \theta | y_1, \dots, y_{25})$.

Use the following "semi-conjugate" prior defined by

$$p(\mu, \theta) = p(\mu)p(\theta)$$

$$p(\mu) = \text{Gaussian}(\mu | \varphi_1, \Phi_1)$$

$$p(\theta) = \text{Wishart}(\theta | \varphi_2, \Phi_2)$$

You will need to figure out some reasonable values of φ_1 , φ_2 , Φ_1 , Φ_2 . Note, I recommend setting φ_2 to a value which has the form (positive integer) / 2 so it is easy to simulate from.

Grading

Nominally this homework is worth 5 points. To get these 5 points you need to have:

- Plots of the samples from the Gibbs chain (at least 1000 steps) for both variables μ and θ
- Histograms of the μ and θ values from the Gibbs chain (also a histogram of $1/\theta$)
- A bivariate scatter plot of the pairs (μ, θ) from the Gibbs chain.
- The code needs to run and does everything "by hand" (so doesn't call to Stan.jl or anything black box)
- I would also like to see some statements at the end that argue what you think I could have picked for μ and $1/\theta$

Note: I reserve the right to give bonus points to those students who go above and beyond the nominal requirements listed above.

This notebook was generated using [Literate.jl](https://github.com/fredriekre/Literate.jl) (<https://github.com/fredriekre/Literate.jl>).


Quick plot of the data

```
In [3]: fig,ax = subplots(1)

ys = range(extrema(yi)..., length=1000)
ax.hist(yi, bins = 5, density=True, histtype = "step")

kde_yi = kde(yi)
ax.plot(kde_yi.x, kde_yi.density)

ax.set_xlabel("data range")
```



```
Out[3]: PyObject Text(0.5, 24.0, 'data range')
```

From the density plot, it seems that the data is centered at -2020

Set the prior parameter

```
In [4]: yΦ, logpriorθ, logpriorμ = let yi=yi
        d = length(yi[1])

        # prior μ ~ Gaussian(φ1, Phi1)
        # prior θ ~ Wishart(φ2, Phi2)
```

```

yΦ = (
    φ1 = 0.0,
    Phi1 = 1/1000.0^2,
    φ2 = 1.5,
    Phi2 = 5.0,
    yi = yi,
    n = length(yi),
    ȳ = mean(yi),
)

# Note: be sure to reason about what values are reasonable for
# φ1, Phi1, φ2, Phi2 I use the variables from demo.

logpriorθ = θ -> (yΦ.φ2-1)*log(θ) - yΦ.Phi2*θ
logpriorμ = μ -> - (yΦ.Phi1/2) * (μ - yΦ.φ1) ^ 2

yΦ, logpriorθ, logpriorμ
end

```

Out [4]: ((φ₁ = 0.0, Phi₁ = 1.0e-6, φ₂ = 1.5, Phi₂ = 5.0, yi = [-2017.71062334 88373, -2007.2408980724263, -2005.1154014930028, -2008.5117830345864, -2025.5132574522222, -2028.1854456826702, -2010.5135720350138, -2025.3244890587323, -2018.6094268648226, -2020.6825199021844 ... -1996.1439726063504, -2037.4334342224986, -2024.9749373753116, -2021.822162447521, -2024.1926752522218, -2013.991961096987, -2020.5517441255515, -2034.3415531928802, -2006.4170649918367, -2041.705257535142], n = 25, ȳ = -2019.8868150750704), var"#5#7"{NamedTuple{(:φ₁, :Phi₁, :φ₂, :Phi₂, :yi, :n, :ȳ), Tuple{Float64, Float64, Float64, Float64, Array{Float64, 1}, Int64, Float64}}}(φ₁ = 0.0, Phi₁ = 1.0e-6, φ₂ = 1.5, Phi₂ = 5.0, yi = [-2017.7106233488373, -2007.2408980724263, -2005.1154014930028, -2008.5117830345864, -2025.5132574522222, -2028.1854456826702, -2010.5135720350138, -2025.3244890587323, -2018.6094268648226, -2020.6825199021844 ... -1996.1439726063504, -2037.4334342224986, -2024.9749373753116, -2021.822162447521, -2024.1926752522218, -2013.991961096987, -2020.5517441255515, -2034.3415531928802, -2006.4170649918367, -2041.705257535142], n = 25, ȳ = -2019.8868150750704)), var"#6#8"{NamedTuple{(:φ₁, :Phi₁, :φ₂, :Phi₂, :yi, :n, :ȳ), Tuple{Float64, Float64, Float64, Float64, Array{Float64, 1}, Int64, Float64}}}(φ₁ = 0.0, Phi₁ = 1.0e-6, φ₂ = 1.5, Phi₂ = 5.0, yi = [-2017.7106233488373, -2007.2408980724263, -2005.1154014930028, -2008.5117830345864, -2025.5132574522222, -2028.1854456826702, -2010.5135720350138, -2025.3244890587323, -2018.6094268648226, -2020.6825199021844 ... -1996.1439726063504, -2037.4334342224986, -2024.9749373753116, -2021.822162447521, -2024.1926752522218, -2013.991961096987, -2020.5517441255515, -2034.3415531928802, -2006.4170649918367, -2041.705257535142], n = 25, ȳ = -2019.8868150750704)))

In [5]: *### Sample the prior to get an idea of the prior assumptions*

```


fig, ax = subplots(2)

 $\theta_i$  = [Wishart( $y\Phi.\phi_2$ ,  $y\Phi.\Phi_{i2}$ ) for i=1:10000]
 $\theta_s$  = range(minimum( $\theta_i$ ),stop=maximum( $\theta_i$ ),length=1000)
logprior $\theta_s$  = logprior $\theta$ .( $\theta_s$ )
logprior $\theta_s$  .-= maximum(logprior $\theta_s$ )
prior $\theta_s$  = exp.(logprior $\theta_s$ )
prior $\theta_s$  ./= sum(prior $\theta_s$  .* ( $\theta_s[2]-\theta_s[1]$ ))
ax[1].plot( $\theta_s$ , prior $\theta_s$ )
ax[1].hist( $\theta_i$ , bins = 20, density=true, histtype = "step")
ax[1].set_xlabel(L"$\theta$ values")

 $\mu_i$  = [Gaussian( $y\Phi.\phi_1$ ,  $y\Phi.\Phi_{i1}$ ) for i=1:10000]
 $\mu_s$  = range(minimum( $\mu_i$ ),stop=maximum( $\mu_i$ ),length=10000)
logprior $\mu_s$  = logprior $\mu$ .( $\mu_s$ )
logprior $\mu_s$  .-= maximum(logprior $\mu_s$ )
prior $\mu_s$  = exp.(logprior $\mu_s$ )
prior $\mu_s$  ./= sum(prior $\mu_s$  .* ( $\mu_s[2]-\mu_s[1]$ ))
ax[2].plot( $\mu_s$ , prior $\mu_s$ )
ax[2].hist( $\mu_i$ , bins = 20, density=true, histtype = "step")
ax[2].set_xlabel(L"$\mu$ values")

fig.tight_layout()

```



gibbs functions

In [6]:

```

function simu(θ, yΦ)
     $\bar{y}$ , n      = yΦ. $\bar{y}$ , yΦ.n
     $\phi_1$ , Phi1 = yΦ. $\phi_1$ , yΦ.Phi1

     $\theta_1$  = Phi1 + n *  $\theta$ 
     $u_1$  = ( $\phi_1$  * Phi1 + n *  $\theta$  *  $\bar{y}$ ) /  $\theta_1$ 

    return Gaussian( $u_1$ ,  $\theta_1$ )
end
#
#
function simθ( $\mu$ , yΦ)
     $y_i$ , n      = yΦ. $y_i$ , yΦ.n
     $\phi_2$ , Phi2 = yΦ. $\phi_2$ , yΦ.Phi2

     $\phi_2$ _new =  $\phi_2$  + n / 2
    Phi2_new = Phi2 + sum(( $y_i$  .-  $\mu$ ) .^ 2) ./ 2
    return Wishart( $\phi_2$ _new, Phi2_new)
end
#
#
function gibbs(yΦ;  $\mu$ _init = 0.0, nsteps=500)
     $\theta$ chain = zeros(nsteps)
     $\mu$ chain = zeros(nsteps)
     $\mu$  =  $\mu$ _init
     $\theta$  = 1.0

    for i=1:nsteps

         $\mu$  = simu( $\theta$ , yΦ)
         $\theta$  = simθ( $\mu$ , yΦ)
         $\mu$ chain[i] =  $\mu$ 
         $\theta$ chain[i] =  $\theta$ 
    end
    return  $\theta$ chain,  $\mu$ chain
end
#
#
# Feel free to use copy this structure with your own details filled in

```

Out[6]: gibbs (generic function with 1 method)

Run the gibbs chain

```
In [7]: yΦ = (
    φ1 = 0.0,
    Phi1 = 1/1000.0^2,
    φ2 = 1.5,
    Phi2 = 5.0,
    yi = yi,
    n = length(yi),
    ȳ = mean(yi),
)

θchain, μchain = gibbs(yΦ, μ_init = mean(yi), nsteps = 1000)
```

```
Out [7]: ([0.006162732416414681, 0.010796006715661466, 0.006503220156173058, 0.006362393035143663, 0.00824428270073555, 0.009356992049805845, 0.005406922086951407, 0.010691788683191355, 0.012428877006109494, 0.006935977793736439 ... 0.010332941144318281, 0.009416937946965026, 0.006937217893297985, 0.008266802515370025, 0.01138683142208521, 0.00584550849762938, 0.010536818923713193, 0.009202361891603058, 0.00767537804585528, 0.009590012755445403], [-2019.964107152187, -2023.1401258978476, -2018.3356960258336, -2018.9478155376378, -2020.0992066584013, -2017.0641099637382, -2016.7105575222079, -2018.6479948948206, -2017.1382495185649, -2017.1463865171097 ... -2022.4570607335936, -2018.5661753122963, -2018.1280060211059, -2024.3627696322656, -2021.7440935446562, -2016.6937593718944, -2018.0732732196043, -2018.4712615811836, -2019.8300338620768, -2018.1101067564166])
```

Plot the results

In [8]: *#Plots of the samples from the Gibbs chain (at least 1000 steps) for b*


```
fig, ax = subplots(2)

burnin = 1
nsteps = length( $\theta$ chain[burnin:end])

ax[1].plot( $\theta$ chain[burnin:end])
ax[1].plot(1:nsteps, fill(1/var( $y\Phi$ .yi),nsteps) )
ax[1].set_ylabel(L"\theta")
ax[1].set_xlabel("chain index")

ax[2].plot( $\mu$ chain[burnin:end])
ax[2].plot(1:nsteps, fill( $y\Phi$ . $\bar{y}$ , nsteps) )
ax[2].set_ylabel(L"\mu")
ax[2].set_xlabel("chain index")

fig.tight_layout()
```



Histograms of the μ and θ values from the Gibbs chain also a histogram of $1/\theta$

```

In [18]: fig, ax = subplots(3)

burnin = 100
steps = length(θchain[burnin:end])


deμ = kde(μchain[burnin:end])
x[2].hist(μchain[burnin:end], bins = 20, density=true, histtype = "step")
x[2].plot(kdeμ.x, kdeμ.density)
x[2].set_xlabel(L"$\mu$ values")

deθ = kde(θchain[burnin:end])
x[1].hist(θchain[burnin:end], bins = 20, density=true, histtype = "step")
x[1].plot(kdeθ.x, kdeθ.density)
x[1].set_xlabel(L"$\theta$ values")

de0 = kde(θchain[burnin:end])
x[3].hist(1 ./ θchain[burnin:end], bins = 20, density=true, histtype = "step")
x[3].set_xlabel(L"$1/\theta$ values")
x[3].set_ylabel("density")

fig.tight_layout()

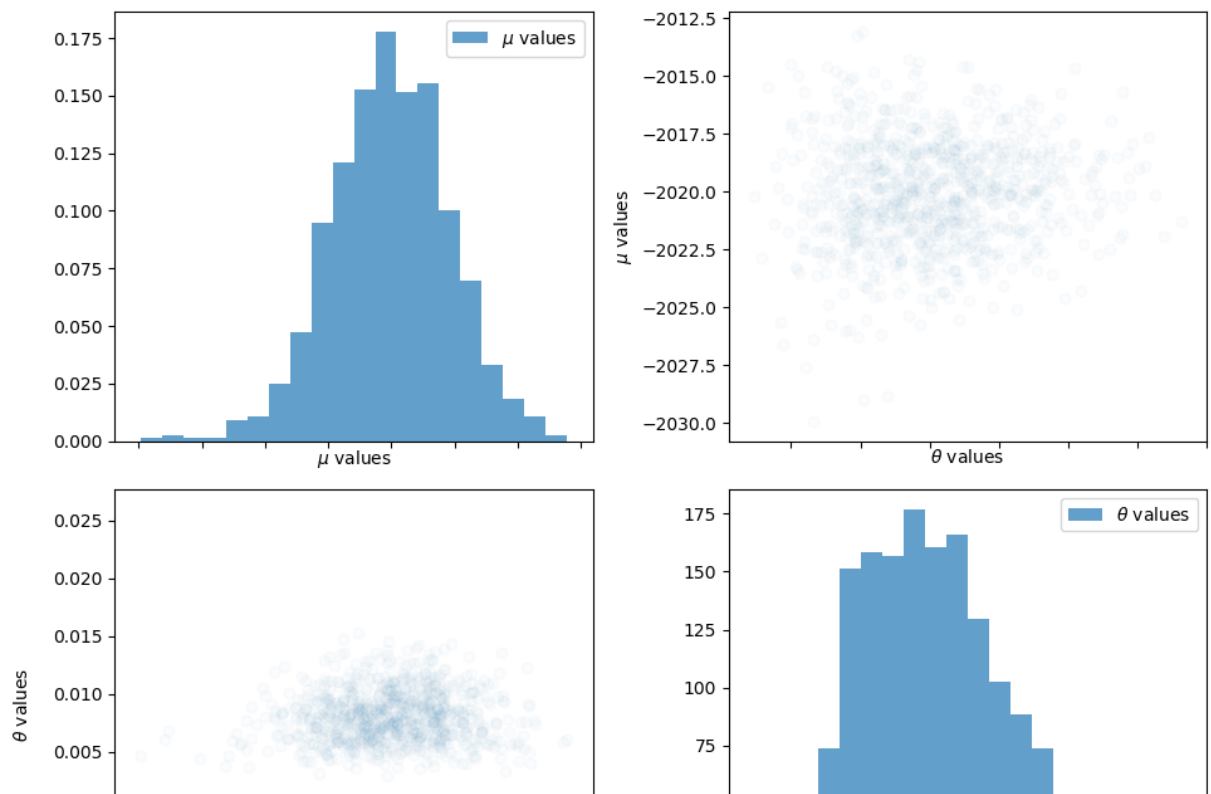
```




A bivariate scatter plot of the pairs (μ, θ) from the Gibbs chain.

```
In [10]: fig, ax = plt.subplots(2, 2, figsize=(10, 8), sharex="col") # , sharey="row"
burnin = 100
nsteps = length(θchain[burnin:end])
# column 1: βi on the xaxis
ax[1,1].set_xlabel(L"$\mu$ values")
ax[2,1].set_xlabel(L"$\mu$ values")
ax[2,1].set_ylabel(L"$\theta$ values")
ax[1,1].hist(μchain[burnin:end], bins = 20, alpha = 0.7, label=L"$\mu$")
ax[2,1].scatter(μchain[burnin:end], θchain[burnin:end], alpha = 0.02)
# column 2: βj on the xaxis
ax[2,2].set_xlabel(L"$\theta$ values")
ax[1,2].set_xlabel(L"$\theta$ values")
ax[1,2].set_ylabel(L"$\mu$ values")
ax[2,2].hist(θchain[burnin:end], bins = 20, alpha = 0.7, label=L"$\theta$")
# kdeθμ = kde((θchain[burnin:end], μchain[burnin:end]))
# ax[1,2].pcolormesh(kdeθμ.x, kdeθμ.y, kdeθμ.density)
ax[1,2].scatter(θchain[burnin:end], μchain[burnin:end], alpha = 0.02)

ax[1,1].legend()
# ax[1,2].legend()
ax[2,2].legend()
#ax[2,1].legend()
fig.tight_layout()
```





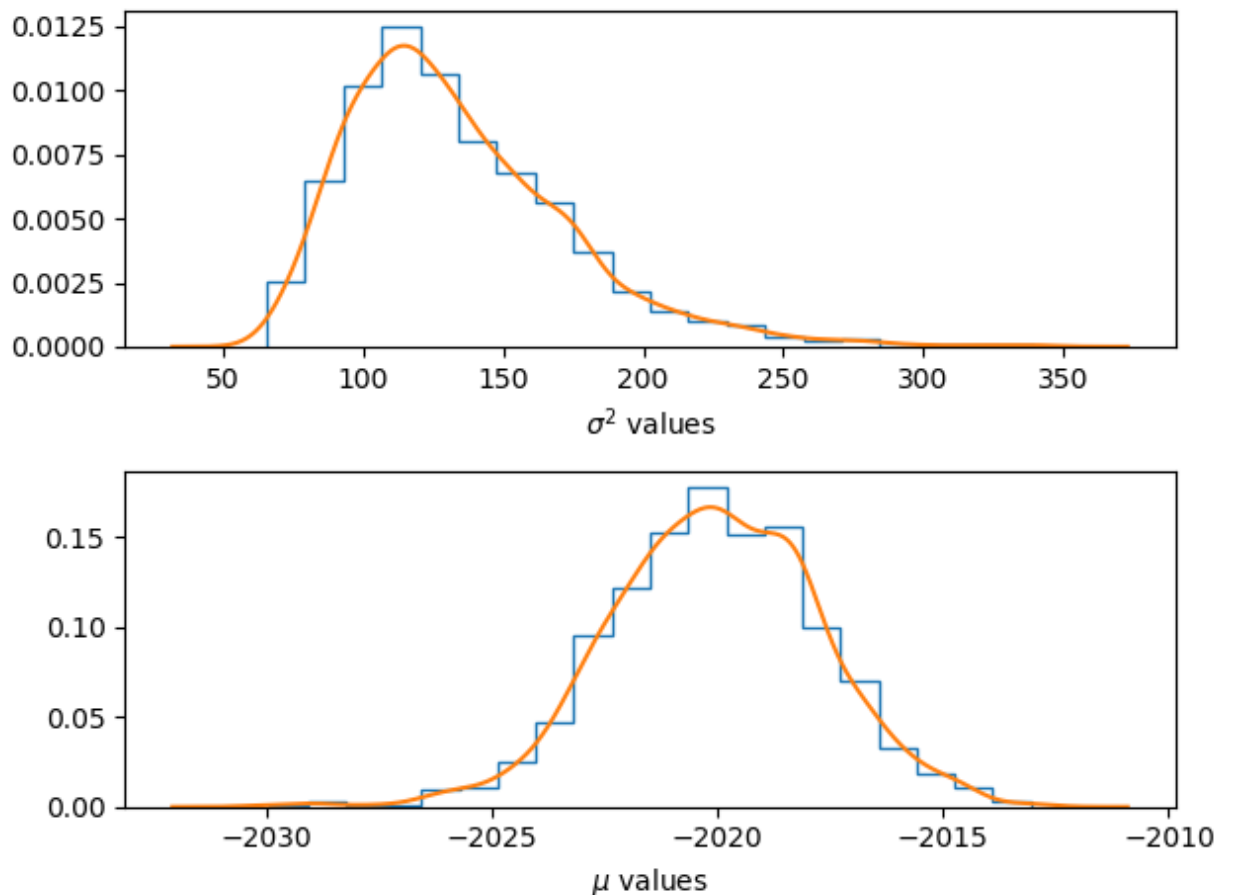
```
In [11]: fig, ax = subplots(2)

burnin = 100
nsteps = length( $\theta$ chain[burnin:end])

kde $\sigma^2$  = kde(1 ./  $\theta$ chain[burnin:end])
ax[1].hist(1 ./  $\theta$ chain[burnin:end], bins = 20, density=true, histtype
ax[1].plot(kde $\sigma^2$ .x, kde $\sigma^2$ .density)
ax[1].set_xlabel(L"$\sigma^2$ values")

kde $\mu$  = kde( $\mu$ chain[burnin:end])
ax[2].hist( $\mu$ chain[burnin:end], bins = 20, density=true, histtype = "st
ax[2].plot(kde $\mu$ .x, kde $\mu$ .density)
ax[2].set_xlabel(L"$\mu$ values")

fig.tight_layout()
```



In [12]: *# plot the prior along with a kde estimate of the posterior*

```
##let  $y\phi = y\phi$ , chain= $\theta$ chain[100:end], logprior = logprior $\theta$ 
let  $y\phi = y\phi$ , chain= $\mu$ chain[100:end], logprior = logprior $\mu$ 

post = kde(chain)
xs = range(-2030, stop=-2010, length=10_000) # you pick the range

fig,ax = subplots(1)
clr = viridis_colors(;length=10)

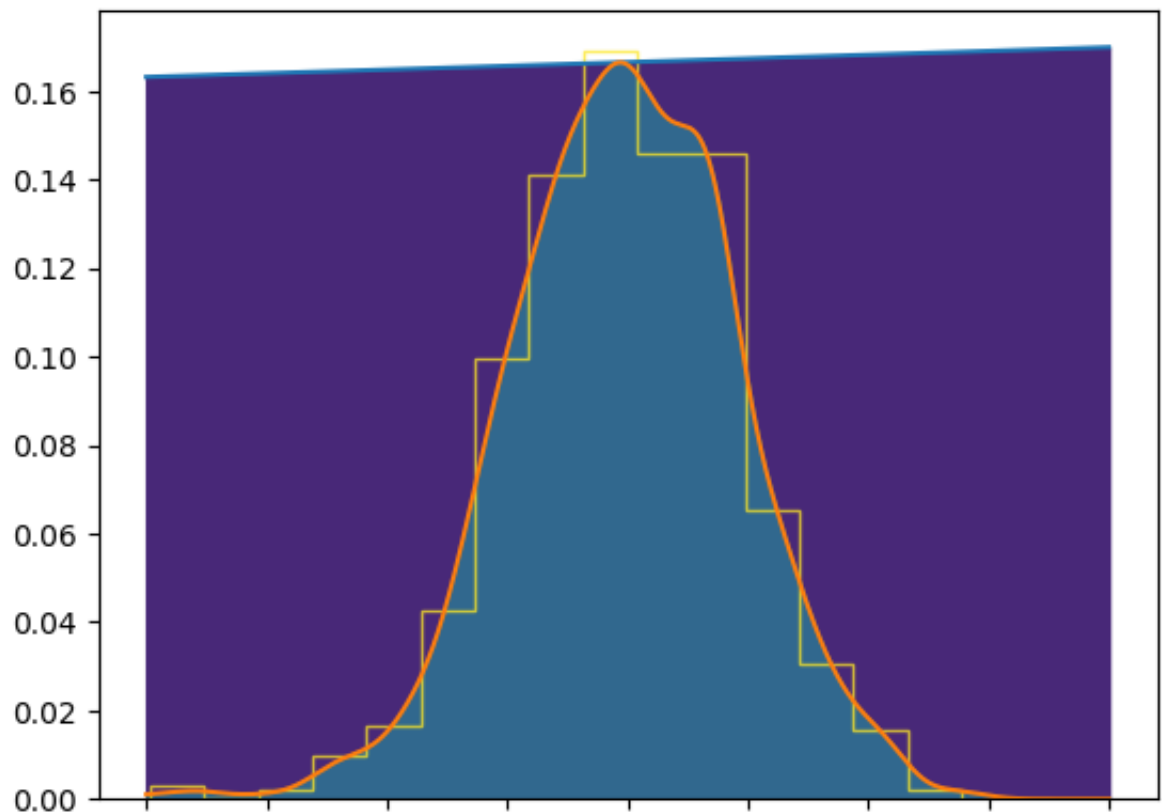
logpriorxs = logprior.(xs)
# -----
logpriorxs .-= logpriorxs[argmax(pdf(post, xs))]
logpriorxs .+= log(maximum(post.density))
priorxs = exp.(logpriorxs)

ax.hist(chain, density=true, histtype="step", bins = 15, color=clr)

ax.plot(xs, priorxs)
#ax.plot(xs, post.density)
ax.plot(xs, pdf(post,xs))

ax.fill_between(xs, priorxs, 0.0, color=clr[2])
ax.fill_between(xs, pdf(post,xs), 0.0, color=clr[4])

end
```



-2030.0-2027.5-2025.0-2022.5-2020.0-2017.5-2015.0-2012.5-2010.0

Out[12]: PyObject <matplotlib.collections.PolyCollection object at 0x157ad4b90>

In [13]: *# plot the prior along with a kde estimate of the posterior*

```

let yΦ = yΦ, chain=θchain[100:end], logprior = logpriorθ

post = kde(chain)
xs = range(0.001, stop=0.018, length=10_000) # you pick the range

fig,ax = subplots(1)
clr = viridis_colors(;length=10)


logpriorxs = logprior.(xs)
# -----
logpriorxs .= logpriorxs[argmax(pdf(post, xs))]
logpriorxs .+= log(maximum(post.density))
priorxs = exp.(logpriorxs)


ax.hist(chain, density=true, histtype="step", bins = 15, color=clr

ax.plot(xs, priorxs)
ax.plot(xs, pdf(post,xs))

ax.fill_between(xs, priorxs, 0.0, color=clr[2])
ax.fill_between(xs, pdf(post,xs), 0.0, color=clr[4])
end

```





Out[13]: Py0bject <matplotlib.collections.PolyCollection object at 0x1590c1690
>

True values of θ and μ

The emperical density of μ peaks at $\mu = -2020$ and the emperical density of peaks at $1/\theta = 120$ (approximately from graph).

If we take the emperical densities for $1/\theta$ and μ as approximations to their true densities, then use the maximum a posteriori (MAP) estimation, then we would estimate $1/\theta = 120$ and $\mu = -2020$.